

# Critical Analysis: Quantum Speedup Claims in Coalition Formation

Technical Report on Mohseni et al. (2024) vs. Rotation Optimization Benchmarks

OQI-UC002-DWave Project  
Technical Analysis Report

December 10, 2025

## Abstract

This report provides a comprehensive technical analysis of the methodology used by Mohseni et al. (arXiv:2405.11917) to demonstrate “quantum scaling advantage” in approximate optimization for coalition formation with 100+ agents. We identify **five critical methodological differences** that explain their claimed 100% solution quality and favorable runtime scaling compared to classical solvers, contrasting with our rotation optimization benchmarks where direct QPU shows 87% optimality gap. Our analysis reveals that their approach uses problem decomposition, specialized hardware (DWaveCliqueSampler), and fundamentally different problem characteristics that dramatically reduce embedding overhead. **Key finding: The speedup is real but highly problem-specific and relies on decomposition + clique embedding, not raw QPU superiority.**

## Contents

# 1 Executive Summary

## 1.1 The Fundamental Question

Why does Mohseni et al. achieve 100% solution quality with D-Wave while our rotation benchmarks show 87% optimality gap?

## 1.2 Key Findings

1. **Problem Decomposition:** They solve many *small* ( $n \leq 20$  variables) independent QUBOs, not one large problem
2. **Clique Embedding:** Uses `DWaveCliqueSampler` which exploits hardware cliques (16-20 qubits) for *zero* embedding overhead
3. **Problem Structure:** Coalition splitting creates *sparse, balanced* graph bisection problems vs. our *dense, frustrated* rotation coupling
4. **Iterative Refinement:** Hierarchical decomposition allows classical+quantum hybrid approach across levels
5. **Apples vs. Oranges:** Benchmarking “approximate optimization” where *any* feasible solution with value  $\geq$  threshold is acceptable

**Verdict:** The speedup is legitimate but contingent on specific problem properties. Not generalizable to arbitrary combinatorial optimization.

# 2 Problem Formulation Comparison

## 2.1 Our Problem: Multi-Period Rotation Optimization

### 2.1.1 Mathematical Formulation

$$\max \sum_{f,c,t} B_c L_f Y_{f,c,t} + \gamma \sum_{f,c',t} R_{c,c'} L_f Y_{f,c,t-1} Y_{f,c',t} + \text{spatial} + \text{penalties} \quad (1)$$

Characteristics:

- **Variables:** 90-900 binary variables ( $5 \times 6 \times 3$  to  $50 \times 6 \times 3$ )
- **Coupling:** Dense temporal + spatial quadratic terms (1860 interactions for 90 vars)
- **Frustration:** 86% negative synergies (spin-glass structure)
- **Constraints:** Soft penalties in objective (CQM  $\rightarrow$  BQM conversion)
- **Embedding Overhead:** 90 logical  $\rightarrow$  651 physical qubits (**7.2 $\times$** )

## 2.2 Their Problem: Coalition Structure Generation (CSG)

### 2.2.1 Mathematical Formulation

Coalition splitting via graph bisection:

$$\min_{x \in \{0,1\}^n} \sum_{i < j} w_{ij} (x_i x_i + x_j x_j - 2x_i x_j) \quad (2)$$

where  $w_{ij}$  = edge weight between agents  $i, j$  in coalition graph.

Characteristics:

- **Variables:** 5-20 binary variables per subproblem (*never* full  $n = 100+!$ )
- **Coupling:** Graph bisection (sparse, balanced cuts preferred)
- **Frustration:** Moderate (not designed to be frustrated)
- **Constraints:** None (unconstrained QUBO)
- **Embedding Overhead:** **Zero** (fits in hardware cliques)

### 2.2.2 Critical Insight: Hierarchical Decomposition

---

**Algorithm 1** Coalition Formation Algorithm (Mohseni et al.)

---

```

1: Start with  $\mathcal{C} = \{[0, 1, \dots, N - 1]\}$  (all agents in one coalition)
2: for iteration = 1 to  $N$  do
3:    $\mathcal{C}_{\text{new}} \leftarrow \mathcal{C}$ 
4:   for each coalition  $c \in \mathcal{C}$  with  $|c| > 1$  do
5:     // Build QUBO for splitting  $c$  into  $c_1, c_2$ 
6:      $Q \leftarrow \text{BuildBisectionQUBO}(c)$    ← Only  $|c|$  variables!
7:      $(c_1, c_2) \leftarrow \text{QPU.Solve}(Q)$    ← Small problem!
8:     if  $V(c_1) + V(c_2) > V(c)$  then
9:        $\mathcal{C}_{\text{new}} \leftarrow \mathcal{C}_{\text{new}} \setminus \{c\} \cup \{c_1, c_2\}$ 
10:      end if
11:    end for
12:    if  $|\mathcal{C}| == |\mathcal{C}_{\text{new}}|$  then
13:      break // No improvement, converged
14:    end if
15:     $\mathcal{C} \leftarrow \mathcal{C}_{\text{new}}$ 
16:  end for
17: return  $\mathcal{C}$ 

```

---

**Key Observations:**

1. QPU *never* sees a problem with 100+ variables
2. Maximum subproblem size  $\approx 20$  variables (fits hardware cliques perfectly)
3. Hundreds of small independent solves vs. one large monolithic solve
4. Classical orchestration + quantum subroutines

## 3 Critical Methodological Differences

### 3.1 Difference 1: Embedding Strategy

#### 3.1.1 What is DWaveCliquesampler?

The D-Wave Pegasus topology contains hardware *cliques* of size 16-20 qubits that are **fully connected** (all-to-all coupling). DWaveCliquesampler automatically:

- Detects if problem fits in a clique ( $n \leq 16$ )
- Maps directly to hardware qubits (bijection, no chains)
- Eliminates embedding overhead entirely

Aspect	Our Approach	Mohseni et al.
Sampler	DWaveSampler + EmbeddingComposite	DWaveCliqueSampler
Logical Variables	90-900	5-20 per solve
Physical Qubits	651 ( $7.2\times$ overhead)	16-20 ( $1.0\times$ overhead)
Embedding Time	43-80 seconds	Negligible (< 1s)
Max Chain Length	9	1 (no chains!)
Chain Breaks	0.02-0.04%	0% (cliques are fully connected)

Table 1: Embedding Comparison

### Code Evidence:

```
# From Main.ipynb line 78
sampler = DWaveCliqueSampler()
coalitions3 = Dwave.solve(value_agent, edges, timeout[num][idx],
                           Dwave_inf, num, idx, rest_inf_dwave, sampler)

# From utils.py line 55
def solve_with_dwave(Q, n, Dwave_inf, num, idx):
    bqm = dimod.BinaryQuadraticModel.from_qubo(Q)
    S = sampler.sample(bqm, num_reads=100) # Direct clique embedding!
    answer = S.lowest().samples()[0]
    solution = [answer[i] for i in range(n)]
    return solution
```

**This is the smoking gun:** Problems with  $n \leq 16$  have *perfect* embeddings. Our 90-variable problem requires complex chain-based embeddings with  $7\times$  overhead.

## 3.2 Difference 2: Problem Size Distribution

Metric	Our Rotation Problem	Their Coalition Splits
Initial Problem Size	90-900 vars	100 agents (full graph)
Subproblem Size	N/A (monolithic)	5-20 vars per split
Number of QPU Calls	1	~100-300 (iterative)
Total QPU Access Time	0.034s (1 solve)	~10-30s (many solves)
Embedding per Solve	43-80s (hard)	< 0.1s (clique)
Total Wall Time	47s (embed + QPU)	40-60s (many fast solves)

Table 2: Problem Scale Comparison

**Analysis:** Their total QPU time is *higher* than ours, but amortized over hundreds of *trivial* embeddings. Our approach spends 95% of time on one giant embedding.

## 3.3 Difference 3: Problem Hardness

### 3.3.1 Frustration and Landscape Structure

#### Our Problem (Rotation):

- 86% negative synergies → frustrated spin glass
- Deep local minima (Gurobi takes 120s, 77K branch-and-bound nodes)
- Integrality gap > 700%

- QPU finds solution with 87% optimality gap (trapped in local minima)

#### Their Problem (Graph Bisection):

- Balanced cut objective (not highly frustrated)
- Any partition with value above threshold is acceptable (“approximate optimization”)
- Problem structure ensures *many* good solutions exist
- QPU samples multiple solutions, picks best among them

#### 3.3.2 What is “100% Solution Quality”?

From the paper abstract: “*quantum annealing on DWave can achieve solutions of comparable quality to our best classical solver.*”

**Key phrase:** “comparable quality”  $\neq$  “optimal.”

Their benchmark is *not* comparing to Gurobi’s exact optimum, but to:

- Tabu search (heuristic)
- Simulated annealing (heuristic)
- QBSolv (D-Wave’s own classical decomposition + SA)

**Translation:** “100% solution quality” means QPU matches the heuristic solvers, not the global optimum. For easy graph bisection problems, heuristics often find near-optimal solutions, so this comparison is favorable.

### 3.4 Difference 4: Constraint Handling

Aspect	Our Approach	Mohseni et al.
Constraint Type	Hard (CQM)	None (unconstrained QUBO)
CQM $\rightarrow$ BQM Conversion	Yes (penalty method)	N/A
Lagrange Multipliers	50.0 (tuned)	N/A
Variable Expansion	90 $\rightarrow$ 120 BQM vars	No expansion
Feasibility Risk	High (penalties may fail)	None (all solutions valid)

Table 3: Constraint Handling Comparison

**Impact:** Our CQM  $\rightarrow$  BQM conversion adds:

- Extra variables (slack variables for inequalities)
- Penalty terms that compete with objective
- Risk of constraint violations if penalties too weak
- Risk of objective suppression if penalties too strong

Their unconstrained QUBO has *none* of these issues.

## 3.5 Difference 5: Benchmarking Philosophy

### 3.5.1 Their Definition: Approximate Optimization

From the paper: “*Approximate optimization is particularly critical for industrial use cases requiring real-time optimization, where finding high-quality solutions quickly is often more valuable than achieving exact solutions more slowly.*”

#### Benchmark Criteria:

- Find *good* solution within time budget
- “Good” = within 5-10% of best known (not necessarily optimal)
- **Time-to-solution** (wall clock) is the primary metric
- Solution quality is *secondary* (as long as “good enough”)

### 3.5.2 Our Implicit Definition: Exact Optimization

#### Benchmark Criteria:

- Compare to Gurobi’s *optimal* solution (ground truth)
- Report optimality gap (our solution vs. true optimum)
- Measure: 87% gap means we found 13% of optimal objective
- Time is secondary (we measure both embedding + solving)

**Implication:** We are solving a *harder* problem (exact optimization of highly frustrated landscape) while they solve an *easier* problem (approximate optimization of moderate landscape).

## 4 Code-Level Analysis

### 4.1 Their Solving Pipeline

```
# Step 1: Build QUBO for coalition split (small!)
def split(c, edges, timeout, Dwave_inf, num, idx):
    Q = {}
    for i in range(len(c)): # len(c) = 5-20 typically
        for j in range(len(c)):
            if i < j:
                utils.add(Q, i, i, edges[(c[i], c[j])])
                utils.add(Q, j, j, edges[(c[i], c[j])])
                utils.add(Q, i, j, -2*edges[(c[i], c[j])])

# Step 2: Solve with DWaveCliqueSampler (fits in clique!)
solution = utils.solve_with_dwave(Q, len(c), Dwave_inf, num, idx)

# Step 3: Split coalition based on solution
c1 = [c[k] for k in range(len(c)) if solution[k] == 1]
c2 = [c[k] for k in range(len(c)) if solution[k] == 0]
return c1, c2

# Clique sampler call (no embedding needed!)
def solve_with_dwave(Q, n, Dwave_inf, num, idx):
    bqm = dimod.BinaryQuadraticModel.from_qubo(Q)
    S = sampler.sample(bqm, num_reads=100) # sampler = DWaveCliqueSampler
    answer = S.lowest().samples()[0]
```

```

solution = [answer[i] for i in range(n)]
return solution

```

Listing 1: Mohseni et al. Solving Pipeline

### Key Observations:

1. QUBO size  $n \times n$  where  $n = |c| \leq 20$
2. Max  $20 \times 20 = 400$  entries, but sparse (only edges present)
3. DWaveCliqueSampler handles  $n \leq 16$  natively
4. For  $n = 17 - 20$ , uses minimal chains (max chain length 2-3)
5. No CQM, no constraint conversion, direct QUBO  $\rightarrow$  QPU

## 4.2 Our Solving Pipeline

```

# Step 1: Build CQM with constraints (large, coupled)
def build_rotation_cqm(data, n_periods=3):
    cqm = ConstrainedQuadraticModel()
    Y = {}
    for f in farm_names: # 5-50 farms
        for c in families_list: # 6 families
            for t in range(1, n_periods + 1): # 3 periods
                Y[(f, c, t)] = Binary(f"Y_{f}_{c}_{t}")

    # Objective: linear + quadratic rotation + spatial (1860 terms!)
    objective = ... # Dense coupling
    cqm.set_objective(-objective)

    # Constraints: <= 2 crops per period per farm
    for f in farm_names:
        for t in range(1, n_periods + 1):
            cqm.add_constraint(sum(Y[(f, c, t)] for c in families_list) <= 2)

    return cqm # 90 vars, 15 constraints

# Step 2: Convert CQM to BQM (penalty method, adds variables)
bqm, info = cqm_to_bqm(cqm, lagrange_multiplier=50.0)
# Result: 90 logical -> 120 BQM variables, 1860 quadratic terms

# Step 3: Find embedding (expensive!)
embedding = find_embedding(source, target, timeout=200)
# Result: 120 BQM vars -> 651 physical qubits (takes 43-80s)

# Step 4: Sample on QPU
sampler = FixedEmbeddingComposite(qpu, embedding)
sampleset = sampler.sample(bqm, num_reads=1000)
# Result: 0.034s QPU time, but 87% optimality gap (trapped in local minima)

```

Listing 2: Our Rotation Solving Pipeline

### Key Differences:

- CQM with constraints (they have none)
- 90 logical vars vs. their  $\leq 20$
- Penalty conversion increases to 120 BQM vars
- Embedding 120 vars to 651 qubits ( $7\times$  overhead)

- Dense coupling (1860 interactions) vs. their sparse bisection
- 86% frustration vs. their balanced cuts

## 5 Why the Speedup is Real (But Limited)

### 5.1 Legitimate Advantages

#### 1. Clique Embedding Eliminates Overhead

- For  $n \leq 16$ , embedding is *free* (direct mapping)
- No chains → no chain breaks → no error accumulation
- Scales to hundreds of independent small solves efficiently

#### 2. Decomposition Matches Hardware Constraints

- Coalition formation is *naturally* hierarchical
- Each subproblem is independent (parallel QPU calls)
- Problem structure “fits” the hardware topology

#### 3. Approximate Optimization Lowers Bar

- Don’t need exact optimum, just “good enough”
- Quantum annealing finds many diverse solutions quickly
- Classical refinement can improve QPU output

### 5.2 Limitations and Non-Generalizability

#### 1. Problem Size Limitation

- Only works for problems that decompose into  $n \leq 16$  subproblems
- For  $n > 20$ , embedding overhead reappears
- Cannot handle monolithic problems (like our rotation)

#### 2. Problem Structure Requirement

- Requires natural decomposition strategy (not always available)
- Graph bisection is “easy” for quantum annealers (balanced cuts)
- Frustrated, dense problems don’t decompose well

#### 3. Comparison Bias

- Benchmarked against heuristics (Tabu, SA), not exact solvers
- Gurobi mentioned but not used as ground truth
- “100% solution quality” relative to other heuristics, not optimum

## 6 Apples-to-Apples Comparison: What If We Used Their Method?

### 6.1 Hypothetical: Decompose Rotation Problem

Could we decompose our rotation problem to fit cliques?

#### Option 1: Decompose by Farm

- Each farm:  $6 \text{ families} \times 3 \text{ periods} = 18 \text{ variables}$
- **Fits in clique!** ( $n = 18 \leq 20$ )
- **Problem:** Rotation synergies couple farms *temporally* (not just spatially)
- **Problem:** Spatial coupling between farms requires coordination

#### Option 2: Decompose by Period

- Each period:  $5 \text{ farms} \times 6 \text{ families} = 30 \text{ variables}$
- **Too large for clique!** ( $n = 30 > 20$ )
- **Problem:** Rotation synergies couple periods (can't separate)

**Verdict:** Our problem has *dense global coupling* that resists decomposition. Their problem has *local structure* amenable to hierarchical splitting.

### 6.2 Hypothetical: Solve Coalition Formation with Our Method

What if they used `EmbeddingComposite` instead of `DWaveCliqueSampler`?

- For  $n = 20$  variables:  $\sim 50\text{-}100$  physical qubits, embedding time  $\sim 1\text{-}5\text{s}$
- Still manageable, but  $10\text{-}50\times$  slower embedding per solve
- Over 300 solves:  $300\text{-}1500\text{s}$  embedding overhead (vs. their  $<10\text{s}$ )
- **Speedup disappears!** Classical would win.

**Conclusion:** Their speedup is *contingent* on clique embedding availability.

## 7 Recommendations for Our Project

### 7.1 What We Should *Not* Do

1. **Don't expect direct QPU to match Gurobi on rotation:** Our 87% gap is realistic for this problem class
2. **Don't use EmbeddingComposite for large monolithic problems:** Embedding overhead destroys quantum advantage
3. **Don't compare approximate to exact optimization:** Different goals  $\rightarrow$  different metrics

## 7.2 What We *Should* Do

### 1. Implement Hierarchical Decomposition:

- Decompose rotation by period (solve periods sequentially)
- Use classical coordination between periods
- Decompose spatial coupling via Louvain/Spectral clustering
- Target subproblem size  $\leq 20$  variables

### 2. Explore DWaveCliqueSampler:

- Redesign formulation to fit cliques ( $n \leq 16$ )
- Use iterative refinement (solve small subproblems repeatedly)
- Accept approximate solutions (don't aim for exact optimum)

### 3. Use Hybrid Solvers Properly:

- LeapHybridCQMSampler for full rotation problem
- Let D-Wave's classical decomposition handle orchestration
- Focus on time-to-good-solution, not time-to-optimal

### 4. Benchmark Fairly:

- If using approximate optimization, compare to heuristics (SA, Tabu)
- If using exact optimization, compare to Gurobi
- Report both optimality gap *and* time-to-solution
- Be transparent about which regime we're in

## 8 Conclusion

### 8.1 Summary of Findings

Factor	Mohseni et al. (Speedup)	Our Project (No Speedup)
Problem Size	5-20 vars per subproblem	90-900 vars (monolithic)
Embedding	DWaveCliqueSampler (zero overhead)	EmbeddingComposite ( $7\times$ overhead)
Problem Structure	Sparse graph bisection	Dense temporal + spatial coupling
Frustration	Low (balanced cuts)	High (86% negative synergies)
Constraints	None (unconstrained QUBO)	Hard constraints (CQM $\rightarrow$ BQM)
Decomposition	Natural hierarchical	Resists decomposition
Benchmark Goal	Approximate (match heuristics)	Exact (match Gurobi)
Solution Quality	100% (vs. heuristics)	13% (vs. optimal)

Table 4: Comprehensive Comparison

## 8.2 Final Verdict

Is their speedup real? **Yes**, for problems with:

1. Natural decomposition into  $n \leq 16$  subproblems
2. Sparse, balanced structure (not highly frustrated)
3. No hard constraints
4. Approximate optimization acceptable

Is their speedup generalizable? **No**, because:

1. Clique embedding only works for tiny problems ( $n \leq 16$ )
2. Decomposition strategy is problem-specific
3. Highly coupled, frustrated problems cannot be decomposed
4. Monolithic large problems hit embedding wall

## 8.3 Implications for Quantum Advantage

- **Quantum advantage exists**, but in a *narrow regime*
- Success requires **algorithm-hardware co-design**
- Problem must be **reformulated** to exploit hardware topology
- Direct QPU is **not** a drop-in replacement for classical solvers
- Hybrid (classical + quantum) is the practical path forward

**Key Takeaway:** Mohseni et al. demonstrate speedup by *engineering the problem to fit the hardware*, not by raw quantum superiority. This is legitimate but requires careful problem selection and reformulation. Our rotation problem, as currently formulated, is fundamentally incompatible with their approach.

## 8.4 Recommended Next Steps

1. **Immediate:** Test hierarchical decomposition (period-by-period)
2. **Short-term:** Benchmark `LeapHybridCQMSampler` vs. Gurobi
3. **Medium-term:** Redesign rotation formulation for clique compatibility
4. **Long-term:** Develop problem-specific decomposition strategies

*“The quantum advantage is not a matter of hardware speed, but of algorithm-problem-hardware alignment.” — This analysis (2025)*