

Benchmark Scenario Analysis: Understanding Test Configurations Across Experiments Multi-Period Crop Rotation Optimization

OQI-UC002-DWave Project

December 2025

Abstract

This report provides a comprehensive analysis of the different benchmark scenarios used across our quantum advantage experiments. We identify four distinct experimental configurations, each designed to test specific aspects of quantum vs classical performance. Understanding these differences is critical for interpreting results and comparing performance across experiments.

Contents

1 Executive Summary	3
2 Benchmark Configuration Taxonomy	3
2.1 Overview Table	3
3 Configuration A: Statistical Comparison (Small Scale)	3
3.1 Script: <code>statistical_comparison_test.py</code>	3
3.1.1 Design Goals	3
3.1.2 Configuration Details	4
3.1.3 Key Characteristics	4
3.1.4 Why This Configuration?	4
4 Configuration B: Hierarchical Statistical (Large Scale)	5
4.1 Script: <code>hierarchical_statistical_test.py</code>	5
4.1.1 Design Goals	5
4.1.2 Configuration Details	5
4.1.3 Three-Level Hierarchical Approach	5
4.1.4 Key Differences from Configuration A	6
4.1.5 Critical Design Choice	6
5 Configuration C: QPU Benchmark (Comprehensive)	6
5.1 Script: <code>qpu_benchmark.py</code>	6
5.1.1 Design Goals	6
5.1.2 Configuration Details	6
5.1.3 Methods Tested	7
5.1.4 Scenarios Tested	7
5.1.5 Why This Configuration?	7

6 Configuration D: Significant Scenarios	7
6.1 Script: <code>significant_scenarios_benchmark.py</code>	7
6.1.1 Design Goals	7
6.1.2 Six Significant Scenarios	8
6.1.3 Configuration Details	8
6.1.4 Methodological Switch	8
6.1.5 Key Difference: Shorter Timeout	8
7 Cross-Configuration Comparison	9
7.1 Timeout Comparison	9
7.2 Formulation Comparison	9
7.3 Method Comparison	9
7.4 When Results Are Comparable	9
8 Interpreting the Runtime Plot	10
8.1 Expected Patterns in All-Runtime-Traces Plot	10
8.1.1 Vertical Groupings	10
8.1.2 Horizontal Divergence	10
8.1.3 Method Clustering	11
8.2 Reading the Plot Correctly	11
9 Recommendations for Future Benchmarks	11
9.1 Standardization	11
9.2 Metadata Requirements	12
9.3 Plotting Guidelines	12
10 Conclusion	12

1 Executive Summary

Our benchmark suite comprises **four distinct experimental configurations**, each with different:

- Problem formulations (6 families native vs 27 foods aggregated)
- Solver methods (clique decomposition vs spatial-temporal vs hierarchical)
- Timeout settings (100s vs 300s vs 900s)
- Problem sizes (4-225 farms, 20-18,225 variables)

Key Finding: Results are **not directly comparable** across all configurations due to these fundamental differences. Within-configuration comparisons are valid; cross-configuration comparisons require careful interpretation.

2 Benchmark Configuration Taxonomy

2.1 Overview Table

Config	Script	Farms	Foods	Timeout	Method
A. Statistical (Small)	statistical_comparison	5-25	6 native	300s	Clique, Spatial-Temporal
B. Hierarchical (Large)	hierarchical_selection	25-100	27(o)6	300s	Hierarchical
C. QPU Benchmark	qpu_benchmark	1-500	5-27	200s	Multiple methods
D. Significant Scenarios	significant_scenarios	5-100	6 or 27	100s	Clique or Hierarchical

3 Configuration A: Statistical Comparison (Small Scale)

3.1 Script: `statistical_comparison_test.py`

3.1.1 Design Goals

- Rigorous statistical comparison on small-medium problems
- Test two QPU decomposition strategies
- Generate publication-quality plots
- Establish baseline quantum advantage metrics

3.1.2 Configuration Details

Parameter	Value
Farm sizes	5, 10, 15, 20, 25
Crop representation	6 families (native)
Periods	3
Variables	90, 180, 270, 360, 450
Gurobi timeout	300 seconds (5 minutes)
QPU reads	100 per subproblem
Iterations	3 (boundary coordination)
Runs per method	2 (statistical variance)

QPU Methods:	
Clique Decomposition	Per-farm decomposition (18 vars/farm)
Spatial-Temporal	Cluster + time decomposition

3.1.3 Key Characteristics

Formulation: Native 6-family representation

- Families: Legumes, Grains, Vegetables, Roots, Fruits, Other
- Each farm-period: one family assignment
- Variables: $F \times 6 \times 3$ where F = number of farms
- Subproblem size: 18 variables (fits in Pegasus clique)

Solver Strategy:

1. **Clique Decomposition:** Solve each farm independently across all periods using DWave-CliqueSampler (native embedding, zero overhead)
2. **Spatial-Temporal:** Decompose by both space (farm clusters) and time (periods), solve iteratively

Post-Processing:

- Optional: Refine family assignments to specific crops
- 3 crops per family
- Diversity analysis (Shannon entropy, unique crops)

3.1.4 Why This Configuration?

- **Clique-friendly:** 18 variables per subproblem fits perfectly in Pegasus cliques (15-20 qubits)
- **Direct comparison:** Same problem structure as Phase 2 roadmap tests
- **Statistical rigor:** Multiple runs enable variance analysis
- **Realistic scale:** 5-25 farms represents typical regional optimization

4 Configuration B: Hierarchical Statistical (Large Scale)

4.1 Script: hierarchical_statistical_test.py

4.1.1 Design Goals

- Scale beyond Configuration A (25-100 farms)
- Test hierarchical decomposition with aggregation
- Demonstrate quantum advantage on larger problems
- Compare against same Gurobi timeout (300s) as Config A

4.1.2 Configuration Details

Parameter	Value
Farm sizes	25, 50, 100
Crop representation	27 foods (o) 6 families (aggregated)
Periods	3
Variables (original)	2025, 4050, 8100
Variables (aggregated)	450, 900, 1800
Gurobi timeout	300 seconds (SAME as Config A)
QPU reads	100 per cluster
Iterations	3 (boundary coordination)
Runs per method	2 (statistical variance)
Farms per cluster	5 (creates 90-var clusters)

4.1.3 Three-Level Hierarchical Approach

Level 1: Aggregation + Decomposition

- Aggregate 27 foods to 6 families (4.5(imes) reduction)
- Partition farms into spatial clusters (5 farms each)
- Cluster variables: $5 \times 6 \times 3 = 90$ (matches Config A scale!)

Level 2: QPU Solving

- Solve each cluster on D-Wave QPU (90 vars fits in clique)
- 3 iterations with boundary coordination
- Total QPU calls: $(F/5) \times 3$ iterations

Level 3: Post-Processing

- Refine family assignments to specific foods
- Local optimization within each family (3-5 foods)
- Diversity analysis and constraint verification

4.1.4 Key Differences from Configuration A

Aspect	Config A	Config B
Farm sizes	5-25	25-100
Crop model	6 families (native)	27 foods (o) 6 families
Variables	90-450	450-1800 (after aggregation)
QPU method	Clique or Spatial-Temporal	Hierarchical only
Cluster size	2-3 farms	5 farms (fixed)
Post-processing	Optional	Required (family(o)food)

4.1.5 Critical Design Choice

Why 5 farms per cluster?

“CRITICAL: Cluster size must create problems comparable to statistical test. Statistical test: 5-25 farms = 90-450 vars (6 families (imes) 3 periods). Our clusters: 5 farms (imes) 6 families (imes) 3 periods = 90 vars per cluster (COMPARABLE!)”

This ensures:

- Each cluster has 90 variables (same as 5-farm problem in Config A)
- QPU solving time comparable to Config A subproblems
- Fair comparison: same-size subproblems, different decomposition

5 Configuration C: QPU Benchmark (Comprehensive)

5.1 Script: qpu_benchmark.py

5.1.1 Design Goals

- Test pure QPU methods (no hybrid solvers)
- Explore multiple decomposition strategies
- Comprehensive timing breakdown
- Wide range of problem sizes (4-500+ farms)

5.1.2 Configuration Details

Parameter	Value
Farm sizes	4-500+ (flexible)
Crop representation	5-27 foods (scenario-dependent)
Variables	20-13,500+
Timeout	200 seconds (direct QPU)
QPU reads	100, 500, 1000, 5000 (configurable)
Annealing time	20, 100, 200 μ s (configurable)

5.1.3 Methods Tested

1. **Direct QPU:** DWaveSampler + EmbeddingComposite (full problem)
2. **Clique Sampler:** DWaveCliqueSampler (for small problems (eq)16 vars)
3. **Manual Decomposition Methods:**
 - PlotBased (per-farm decomposition)
 - Multilevel (hierarchical partitioning)
 - Louvain (community detection)
 - Cutset (graph cuts)
 - Spectral (spectral clustering)

5.1.4 Scenarios Tested

Synthetic Scenarios (simple assignment, no rotation):

- micro_6 (2 farms, 3 foods, 6 vars)
- tiny_24 (4 farms, 6 foods, 24 vars)
- small_100 (20 farms, 5 foods, 100 vars)
- medium_160 (32 farms, 5 foods, 160 vars)

Rotation Scenarios (3-period rotation with synergies):

- rotation_micro_25 (5 farms, 5 foods, 3 periods, 75 vars)
- rotation_small_50 (10 farms, 5 foods, 3 periods, 150 vars)
- rotation_medium_100 (20 farms, 5 foods, 3 periods, 300 vars)
- rotation_large_200 (40 farms, 5 foods, 3 periods, 600 vars)

5.1.5 Why This Configuration?

- **Comprehensive exploration:** Tests many decomposition strategies
- **Timing granularity:** Separates CQM build, BQM conversion, embedding, solving
- **Flexible scale:** From tiny (6 vars) to massive (13,500+ vars)
- **Pure quantum:** No hybrid solvers, understand true QPU capabilities

6 Configuration D: Significant Scenarios

6.1 Script: significant_scenarios_benchmark.py

6.1.1 Design Goals

- Unified benchmark across all problem sizes
- Test scenario-appropriate methods (clique for small, hierarchical for large)
- Consistent timeout across all sizes
- Direct head-to-head Gurobi vs QPU comparison

6.1.2 Six Significant Scenarios

Name	Farms	Foods	Vars	QPU Method
rotation_micro_25	5	6	90	clique_decomp
rotation_small_50	10	6	180	clique_decomp
rotation_medium_100	20	6	360	clique_decomp
rotation_250farms	25	27	2025	hierarchical
rotation_350farms	50	27	4050	hierarchical
rotation_500farms	100	27	8100	hierarchical

6.1.3 Configuration Details

Parameter	Value
Gurobi timeout	100 seconds (DIFFERENT from A, B)
MIP gap	0.01 (1% optimality)
MIP focus	1 (find feasible quickly)
Improve start time	30s (stop if no improvement)
QPU reads	100
Farms per cluster	5 (hierarchical)
Iterations	3 (boundary coordination)

6.1.4 Methodological Switch

Why different methods for different sizes?

- **Small (5-20 farms, 6 foods):** Use clique_decomp
 - Subproblems fit perfectly in Pegasus cliques
 - Direct QPU embedding, no overhead
 - Proven effective in Config A
- **Large (25-100 farms, 27 foods):** Use hierarchical
 - Too large for direct embedding
 - Requires food(o)family aggregation
 - 3-level decomposition proven in Config B

6.1.5 Key Difference: Shorter Timeout

`GUROBI_CONFIG = {'timeout': 100, ...} ← 100s, not 300s!`

Rationale:

- Emphasis on *practical* time-to-solution
- Real-world constraint: users want results quickly
- Demonstrates quantum advantage more clearly
- Still long enough for Gurobi to find good solutions

7 Cross-Configuration Comparison

7.1 Timeout Comparison

Config	Gurobi Timeout	QPU Timeout	Rationale
A (Statistical)	300s	N/A	Match Phase 2 roadmap
B (Hierarchical)	300s	N/A	Consistency with Config A
C (QPU Benchmark)	120-300s	200s	Comprehensive timing
D (Significant)	100s	N/A	Practical emphasis

Implication: Speedup factors are **not directly comparable** across configurations due to different timeout settings.

7.2 Formulation Comparison

Config	Crop Model	Aggregation	Post-Processing
A	6 families	None	Optional
B	27 foods (o)6	Yes (4.5(imes))	Required
C	5-27 foods	Varies	Optional
D	6 or 27	Size-dependent	Method-dependent

Implication: Objective values are **not directly comparable** between native 6-family and aggregated 27(o)6 formulations.

7.3 Method Comparison

Config	QPU Methods
A	Clique Decomposition, Spatial-Temporal Decomposition
B	Hierarchical Decomposition only
C	Direct QPU, Clique, Multiple Manual Decomposition strategies
D	Clique Decomposition (small) or Hierarchical (large)

7.4 When Results Are Comparable

Valid Within-Configuration Comparisons:

- Config A: Clique vs Spatial-Temporal (same timeout, formulation)
- Config B: Gurobi vs Hierarchical QPU (same timeout, formulation)
- Config C: Different decomposition methods (same scenario)
- Config D: Gurobi vs QPU (same scenario, but note timeout difference)

Invalid Cross-Configuration Comparisons:

- Config A vs Config D: Different timeouts (300s vs 100s)
- Config A vs Config B: Different formulations (native 6 vs aggregated 27(o)6)
- Absolute speedups: Depend on timeout setting

- Absolute gaps: Depend on formulation and timeout

Valid Cross-Configuration Observations:

- Scaling trends (how performance changes with problem size)
- Method effectiveness relative to problem size
- Consistency of quantum advantage (present across all configs)

8 Interpreting the Runtime Plot

8.1 Expected Patterns in All-Runtime-Traces Plot

When viewing the comprehensive runtime plot with all traces:

8.1.1 Vertical Groupings

What you'll see:

- Multiple traces at same variable count with different runtimes
- E.g., at 90 variables: multiple Gurobi traces, multiple QPU traces

Why this happens:

- Different configurations test same problem size
- Config A: 5 farms (5 times) 6 families (5 times) 3 periods = 90 vars
- Config B cluster: 5 farms (5 times) 6 families (5 times) 3 periods = 90 vars
- Config C scenario: varies by specific test
- Different timeout settings (5) different Gurobi runtimes

8.1.2 Horizontal Divergence

Gurobi traces:

- Config A, B: Plateau at 300s (timeout)
- Config D: Plateau at 100s (shorter timeout)
- Config C: Variable timeout (120-300s)

QPU traces:

- Generally increasing with problem size
- Sub-linear scaling for decomposition methods
- Multiple traces for same size (different files/runs)

8.1.3 Method Clustering

Expected clusters:

1. **Gurobi family** (circles, blue shades):
 - ground_truth (Config A)
 - gurobi (Configs B, C, D)
 - gurobi_rotation (Config C roadmap tests)
2. **Clique family** (triangles, red shades):
 - clique_decomp (Configs A, D)
 - clique_qpu (Config C)
 - clique_decomposition (Config C)
3. **Spatial/Hierarchical family** (triangles, orange shades):
 - spatial_temporal (Config A)
 - spatial_temporal_decomp (Config C)
 - hierarchical_qpu (Config B)
 - hierarchical (Configs B, C, D)
4. **Other QPU methods** (squares, green shades):
 - direct_qpu (Config C)
 - quantum (Config C scaling tests)

8.2 Reading the Plot Correctly

Do compare:

- Traces from same file/configuration
- Scaling trends within method families
- General quantum vs classical patterns

Don't compare:

- Absolute runtimes across configurations (different timeouts)
- Speedup values (need consistent baselines)
- Single traces in isolation (need context)

9 Recommendations for Future Benchmarks

9.1 Standardization

To enable direct comparison, future benchmarks should:

1. **Unified timeout:** Use consistent Gurobi timeout (e.g., 300s)
2. **Clear labeling:** Include config identifier in result files
3. **Consistent formulation:** Specify native vs aggregated
4. **Method tagging:** Explicitly tag decomposition strategy

9.2 Metadata Requirements

Every result file should include:

```
{  
    "configuration": "A_statistical_small",  
    "script": "statistical_comparison_test.py",  
    "formulation": "6_families_native",  
    "gurobi_timeout": 300,  
    "timestamp": "2025-12-15T...",  
    "parameters": {...}  
}
```

9.3 Plotting Guidelines

When creating comparison plots:

1. **Filter by configuration:** Plot one config at a time
2. **Annotate differences:** Mark timeout boundaries
3. **Use facets:** Separate panels for different formulations
4. **Color by method:** Consistent color scheme across plots
5. **Legend clarity:** Include config identifier

10 Conclusion

Our benchmark suite comprises four distinct configurations, each designed for specific research questions:

- **Config A:** Small-scale statistical comparison (5-25 farms, clique methods)
- **Config B:** Large-scale hierarchical scaling (25-100 farms, aggregation)
- **Config C:** Comprehensive QPU exploration (4-500+ farms, multiple methods)
- **Config D:** Unified significant scenarios (5-100 farms, adaptive methods)

Key Takeaway: Results are valid within configurations but require careful interpretation across configurations due to:

1. Different timeout settings (100s vs 300s)
2. Different formulations (native 6 vs aggregated 27(o)6)
3. Different QPU methods (clique vs hierarchical)
4. Different problem scales

Consistent Finding: Quantum advantage (speedup >1) is demonstrated across *all* configurations, providing robust evidence of QPU effectiveness for crop rotation optimization.